

# HIGH LEVEL DESIGN

HLD 101

## Difference between HLD & LLD

HLD explains how the relationships between the components of a software system are designed.

LLD talks more about how the code inside these components is structured so that if any change needs to be brought in, minimum lines of code will be changed.

**Vertical Scaling** - When you replace the current hardware with better hardware.

## Delicio.US

- Delicio.us was an extension (before Chrome) that was used to sync bookmarks across multiple machines and browsers using email-password authentication.
- This company was later bought out by Yahoo.

- How do the client-server requests usually handled?

- User enters the web-url on the browser.
- Browser does not understand the url.  
So, it contacts DNS services - which have access to DNS servers.

**DNS Servers :-** Converts url to ip addresses / Finds the ip addresses of URLs/Website

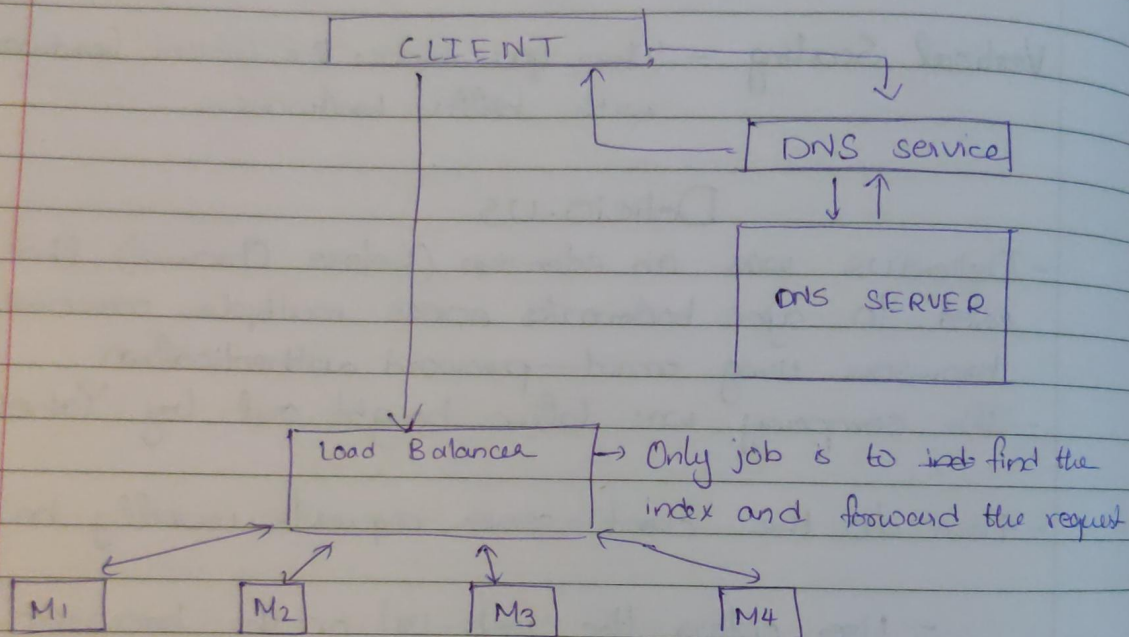
- ICANN maintains the url-ip repos.
- The browser with the IP address can now visit the site on the server.

- Now Delicious.us was hosted on a laptop in the beginning on a post (assume so).
- After a while, the traffic started increasing.

### Vertical Scaling:-

Option 1 - Increase the memory of the current machine.  
Horizontal Scaling  
Option 2 - Use parallel machines.

- for this, you might want to use a load balancer to avoid some of the bottleneck.



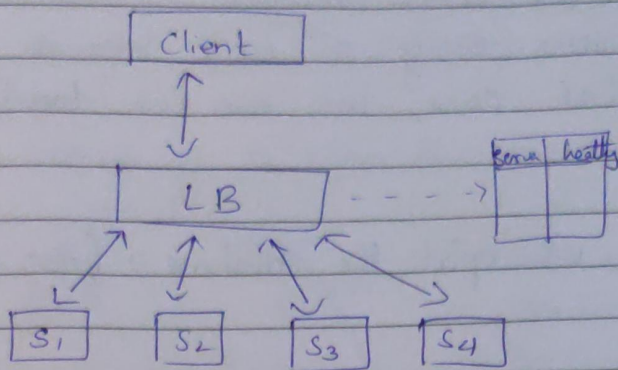
- There are two problems with this.

- ① how does the load balancer know which server to send the request to?
- ② how do we split the DB and store it on multiple machines.



① Let's look into the ways we can utilize a Load Balancer.

- Consider a system (array of servers) only doing a summing  $(a+b)$ .



- To route the requests, we can go with

I. Round Robin

II. Least Response Time

III. Least loaded Server.

IV. Weighted Round Robin

- Consider Round Robin for the time being.

- Here, the problem is that if one server goes down, every 4<sup>th</sup> request is failed.

- We can deal with that issue by checking the health of the server before sending it a request.

- For example, if  $S_3$  is down, we can skip over it and forward ~~it~~ the request to the next healthy server.

- How does the load balancer know if the server is healthy or not? - LB maintains a hashmap.

- This hashmap could be updated using one of the following:

① Heartbeat (pushed by the server - ping once every few moments to let the load balancer know the server is alive.)

② Healthcheck endpoint / url (If the server does not respond in time or if the server sends false or it has reached its capacity, we mark it as unhealthy.)



## - Consider Least Response Time

- Let's say one of the servers is responding very slow for some reason.
- This means every 4<sup>th</sup> user has a bad experience.
- In that case, we can use Least Response Time model.

## ② How do we split the database? (Also called sharding)

### Thoughts :-

#### I. prefix machineID with userID

- In this case, if the current system ran out of memory too and if we want to move the current data to new servers, we would want to change the user IDs which is not a good design. (Cached apps might break).

#### II. userID % no. of servers

- As soon as a new server is added, no. of servers i.e., the denominator is changed. So, almost all the users need to be moved.

## Solution - Consistent Hashing

- Consistent Hashing is used almost everywhere in H10
- Constraints we have
  - Load Balancer cannot have a lot of data
  - Load Balancer should somehow know which server to route the request to
  - Addition or deletion of servers should not cause a lot of data to move.
  - Cannot change userID

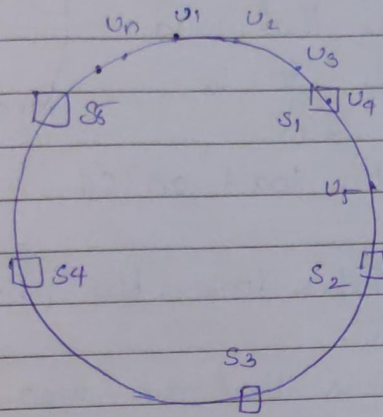
- We will have 2 hash functions

Hash-1 (user\_id)

Takes the user\_id and converts/generates a long long int in the range of  $[1, 10^8]$  inclusive

Hash-2 (server\_id)

Takes the server\_id & does the same.



- When we get a user\_id, we generate hash-1 & assign it to its next-nearest server on the cyclic edge.

- Implementation

Arrange servers in a sorted order, sorted based on their hash values.

server hashes =

9	2800	4000	4300	98000
0 [S1]	1 [S3]	2 [S2]	3 [S4]	4 [S5]

we

- When we get a user\_id, we run hash-1, get the value,  $x$  & perform binary search on "server-hashes" to find the next-greatest server and assign/route the request there.

Addition of the servers:-


- Let's say we added  $S_x$  server and it falls between  $S_3$  &  $S_4$ .
- Now, we can <sup>move</sup> copy all the users that need to be present in  $S_x$  from  $S_4$ . Here, only some data is copied over. So, that is acceptable.



- The concern with this addition is that, its taking load off of only one server.

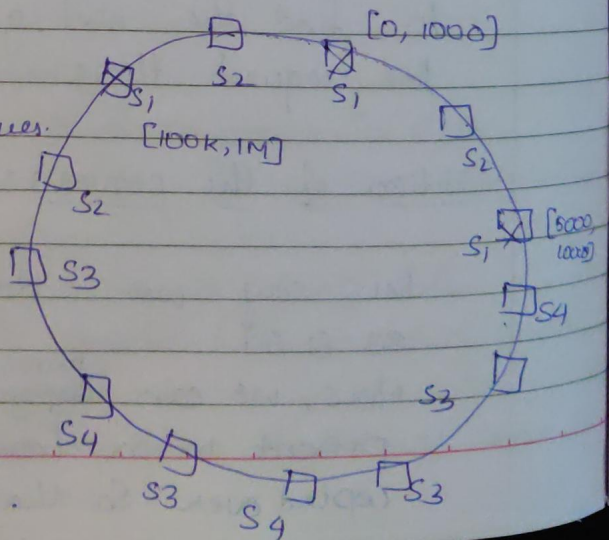
### Deletion of Servers

- Let's assume that for some reason  $s_3$  failed.
- Now, all the users that need to be added/routed to  $s_3$  are diverted to  $s_4$ .
- Now, effectively, the load on  $s_4$  is doubled ( $2x$ ).
- Because of this extra load, the chances are  $s_4$  goes down as well.
- Now all the users that are meant for  $s_3$  &  $s_4$  are diverted to  $s_5$ . Here, the load is  $3x$  now.
- This will probably fail too.
- This phenomenon is called **CASCADING EFFECT FAILURE**.
- To overcome the problem, we go with Modified Consistent hashing.



### Modified Consistent Hashing

- We provide each server with multiple hash values. [look at the figure]
- Let's say  $s_1$  died.
- All the users who match with  $[0, 1000]$  hash will be moved to  $s_2$ .
- $[5k - 10k]$  to  $s_4$ .
- $[100k - 1M]$  to  $s_2$  again.



- ① Since the hash values are spread out randomly b/w  $[1, 10^{18}]$  or some such, we can ensure that all the servers are getting about equal load.

**NOTE :** We do not have 3 physical servers representing each server. We just assign multiple hash values by using multiple hash fns.

server-id  $\longrightarrow$  passes through 5 hash fns (or avg)  
 user-id  $\longrightarrow$  passes through 1 hash fn

- Having more hash fns spread out the data more evenly.
- But this will increase the amount of data stored on the load balancer.

- ② When a new server  $s_x$  is added, it also has multiple hash fns and it will be plotted sparsely all around the cyclic edge.  
 So, the server takes off load from multiple (almost all) servers evenly.

### Hypothetical Hash

- To make it fall in the range  $[1, 10^{18}]$

hash(user-id):

$m = \text{md5}(\text{user-id}); \text{num} = 1$

for  $c$  in  $m$ :

$\text{num} = \text{num} * 26 + c$

$\text{num} = \text{num} \% 10^{18}$

return num



Hash Conflict :-

- What happens if two hashes collide?

S1 →	100	200	500	800	1000
S2 →	80	100	300	900	12000

- You can have a priority. Either  $S1 > S2$  or  $hash_1 > hash_2$  or vice versa.

Stateless LB :- It does not matter which server is contacted / requested.

- Round Robin with health check could be used.

Stateful LB :- It matters which server is requested for.

- When data is sharded.